


Screen 1: Landing Page



Screen 2: Login Page

SiBoNi

Username

Password

[Log In](#)

☐ Remember me

[Forgot Your Password?](#) [Use Custom Domain](#)

or

☒ [Log In with Email](#)

Not a customer? [Try for Free](#)

SiBoNi CXO Cockpit
for executive
decision-making

Read more

Link back to landing page of
an asset or a CTA that we
want existing or prospective
customers to do



Screen 3: Home Page

SiBoNi

User Name, Company Name

Last updated: 2h ago | Sources: Internal systems + external signals | Confidence model: enabled

What changed

External: Steel +6% in 30 days; competitor discounting in OEM bids
Internal: Plano utilization 92%; Axle-B scrap 2.1% → 3.4%
Net: Growth available, capacity + quality are constraints

Top Decision this week

Decision: Unblock 2 late-stage OEM deals (Customer A)
Why now: signature window 10–14 days; discount/re-bid risk rising
Impact: +\$8–\$15M order intake (30–45d) | Confidence: 80%
Next: owner CRO+Legal; approve guardrails; exec sponsor call

Value at stake

Revenue at risk: \$8–\$15M
Margin at risk: 30–70 bps
OTIF/penalty exposure: \$0.8–\$1.6M
Cash if buffering inventory: -\$10M to -\$25M

Click to create a board brief

A. Insights Hub

B. Decision Hub

C. Execution Hub


Click to go wherever CXO wants to go deeper in to

Board Brief


Reset

Snap


Screen 3.A: Insights Hub



Screen 3.B: Decision Hub



Screen 3.B: Decision Hub



Screen 3.C: Execution Hub

Decisions you've committed to, and what's happening now

SiBoNi

Reset

Snap

- Home
- Board Brief
- Insights Hub
- Decision Hub
- Execution Hub

- Time range: Last 30
- Plant: Plano, Texas
- Region: USA
- CustSegment: OEM
- Prod line: Car Axle

Decisions	Owner	Status	Due	Value at Stake
D1: Dual source critical component	COO	Stalled		+\$X revenue, +Y bps margin
D2: Quality containment, Product lineAI	CPO	Completed		
D3: Services and retrofit growth play	COO	In Progress		
D4: Unblock two late stage deal	CSO	Stalled		

Summary:

Total Value at Stake: +\$X revenue, +Y bps margin

Active Decisions: 7In Progress: 18, At Risk: 4, Overdue: 2

Revenue growth

Orders / order intake

Profit Margin

Backlog / book-to-bill

Free cash flow

Safety Performance
(LTIFR-Lost Time Injury
Frequency Rate)

Appendix

Data populated view: CEO

SiBoNi

Orion Industrial Systems

What Changed

External: Steel up ~6% in 30 days, competitor discounting reported in OEM bids.

Internal: Plano utilization at 92%, scrap rate for Axle-B up from 2.1% → 3.4%.

Net: growth is available, but **capacity + quality** are now the constraints.

Top Decision This week

Unblock the 2 late-stage OEM deals with Customer A (Plano | Car Axle) **Why now:** signature window closes in **10–14 days**, re-bid/discount risk rising

Expected impact: **+\$8M–\$15M** order intake in 30–45 days

Confidence: Medium-High (80%)

Next action: set owner (**CRO + Legal**), approve pricing guardrails, schedule exec sponsor call

Value at Stake

Revenue at risk: \$8M–\$15M (2 late-stage Customer A deals)

Margin at risk: 30–70 bps (if discounting + expediting increases)

OTIF / penalty exposure: \$0.8M–\$1.6M (late delivery clauses on 3 programs)

Cash impact (if we buffer inventory): -\$10M to -\$25M (short-term working capital)

Insights Hub

Decision Hub

Execution Hub

Board Brief

↺ Reset

📐 Snap

SiBoNi

Reset

Snap

Home

Board Brief

Insights Hub

Decision Hub

Execution Hub

Time range: Last 30

Plant: Plano, Texas

Region: USA

CustSegment: OEM

Prod line: Car Axle

LIVE

Copper (HG=F): \$5.851/lb

US CPI YoY: 2.4%

Crude Oil: \$96.40 per barrel

Decision

D1: Unblock the two late-stage deals with Cust A [**Conf:** High]

D2: Add capacity-temporary second shift [**Conf:** Low]

D3: Run a TQM sprint in plant for Axle-B [**Conf:** Medium]

Signals

S1: OEM "Atlas Motors" new model ramp announced for Q2

S2: On-time delivery drop Supplier X (97% → 91% over 4 wk)

S3: Quality: Scrap rate Axel-B rise from 2.1% → 3.4%

Impact

I1: Order Intake (\$42.6M (+9% vs prior 30D)

I2: Quote to Order Conversion: 24% (target 28%)

I3: On-time Delivery (OTIF): 93% (target 96%)

Risk

R1: late-stage deals stuck in legal. Revenue shifts out of Qtr

R2: backlog increasing inventory, receivables, hurting cash

R3: a critical component lead time spike blocks build schedule.

Revenue growth

Profit Margin


Free cash flow

Orders / order intake


Backlog / book-to-bill

Safety Performance
(LTIFR-Lost Time Injury
Frequency Rate)

Decision Hub



Decision Hub



Execution Hub

SiBoNi

Reset

Snap

Home

Board Brief

Insights Hub

Decision Hub

Execution Hub

Time range: Last 30

Plant: Plano, Texas

Region: USA

CustSegment: OEM

Prod line: Car Axle

Decisions

Owner

Status

Due

Value at Stake

D1: Dual source critical component

COO

Stalled

+\$X revenue, +Y bps margin

D2: Quality containment, Product lineAI

CPO

Completed

D3: Services and retrofit growth play

COO

In Progress

D4: Unblock two late stage deal

SVP-Sales

In progress

\$8M–\$15M (2 late-stage deals)

Summary:

Total Value at Stake: +\$X revenue, +Y bps margin

Active Decisions: 7

In Progress: 18, At Risk: 4, Overdue: 2

Revenue growth

Profit Margin

Free cash flow

Orders / order intake

Backlog / book-to-bill

Safety Performance
(LTIFR-Lost Time Injury
Frequency Rate)

SiBoNi Developer Appendix

Data contract

Taxonomy, sources, schema,
provenance

Decision contract

Decision method, impact,
confidence, actions

Build contract

Architecture, connectors, trust
posture

A1. Prototype scope and build contract

What the developer should treat as mandatory for the first working prototype

Prototype build contract

- Goal: a working prototype that feels real to CXOs and preserves the core decision flow.
- Screens in scope: Home, Insights Hub, Decision Hub, Execution Hub.
- Core objects: Decision, Signal, KPI, Risk, Action, Owner, Status, Due date.
- Non-negotiables: every decision shows Why now, Impact range, Confidence, and Top 3 evidence items.
- Mock data is acceptable, but every record must carry provenance fields: source, date, extraction method.
- Commit/Hold actions must create or update an Action record with owner, due date, status, and value-at-stake.
- Out of scope for prototype: full connector build, full multi-tenant scale, deep BI visualization.

A2. Signal taxonomy and source strategy

How the system classifies signals and what it is allowed to ingest

Signal taxonomy

- Market & Business: demand shifts, pricing pressure, orders and backlog signals.
- Technology & Innovation: automation, AI adoption, product and platform shifts.
- Risk & Reputation: safety, quality, cyber, legal, and brand exposure.
- Regulatory & Policy: compliance changes, tariffs, reporting obligations.
- Benchmarking & Perception: competitor moves, customer sentiment, analyst narratives.
- Workforce Dynamics: attrition, hiring constraints, labor relations, productivity.
- Alignment & Blindspots: gaps vs priorities, initiatives misaligned, missed signals.

Source strategy

- Public sources: filings, annual reports, investor decks, press releases, reputable news, market indicators.
- Internal sources (Phase 2): CRM, ERP, HRIS, OKRs, pipeline views, quality and delivery metrics.
- Trust rule: every signal stores provenance, source, timestamp, and extraction method.
- Internal data quality is assumed imperfect, use quality gates, human review, and confidence weighting.
- Each signal maps to one primary taxonomy category and can carry optional tags and entities.
- Use the selected user context, region, plant, product line, customer segment, to scope what is surfaced.

A3. Data window, retention, and canonical signal schema

The rules that keep signals fresh, trustworthy, and cheap enough to run

30-day window and retention

- Default horizon: rolling 30-day high-fidelity signal window for decision-grade freshness.
- Older data is summarized and clustered into patterns to reduce noise and inference cost.
- Use incremental updates. Do not reprocess the full corpus if only a subset changed.
- Keep final recommendations and evidence packs. Prune intermediate artifacts where possible.
- Noise control: de-duplication, confidence weighting, and anomaly flags.
- The product should always show when the current view was last refreshed.

Canonical signal schema

- SignalID: unique and immutable.
- Core fields: timestamp, source_type, source_ref, extraction_method.
- Classification: category, tags, keywords, entities such as company, plant, region, product line, segment.
- Content: 1–2 line summary plus short evidence snippet.
- Trust: confidence_score, bias_flags, needs_verification.
- Linkage: related_decision_ids, related_kpi_ids, related_risk_ids.
- Audit: created_by, created_at, version.

A4. Decision method and explainability contract

How recommendations are produced and what the UI must expose

Decision method

- Frame the decision as a choice: what must be decided, by when, and why now.
- Generate 2–4 options, including a conservative path and a bolder path.
- For each option show upside, downside, dependencies, leading indicators, and irreversible commitments.
- Quantify impact on 1–2 KPIs as a range: low, likely, high.
- Run an alignment check against current initiatives: aligned, partial, not aligned.
- Recommend one option with confidence and a short pre-mortem risk list.
- Output a 2-week action plan with owner and checkpoint date.

Explainability contract

- Every decision must show Why now, Impact, Confidence, and Top 3 evidence items.
- Evidence can include representative signals and short source excerpts where available.
- Unsupported claims should be blocked or downgraded.
- Confidence should be High, Medium, or Low with one clear reason.
- The UI should preserve provenance in both the Decision Hub and any exported brief.
- Show what would change the recommendation: missing data, weak evidence, or a new signal.

A5. AI/analytics approach and quality system

What happens behind the scenes and how output quality is controlled

AI and analytics approach

- LLMs are the interface and synthesis layer, not the entire solution.
- Use classical methods alongside LLMs: clustering, topic modeling, anomaly detection, forecasting.
- Use knowledge graph or entity linking to connect companies, plants, products, customers, and topics.
- Run a multi-step pipeline: extract signals, detect patterns, draft insights, generate decision options.
- Apply caching and summarization to reduce token and retrieval cost.
- Design model-provider flexibility behind a model gateway.

Quality system

- Measure extraction precision and recall, pattern stability, evidence coverage, bias indicators, and decision usefulness.
- Keep a small gold-set of signals and briefs for regression testing.
- Run automated checks for provenance presence, entity consistency, and contradiction detection.
- Use human-in-the-loop verification for high-impact claims.
- Support a Needs Verification queue and bias alerts dashboard in the admin surface.
- Track before/after behavior when prompts, extraction logic, or models change.

A6. Reference architecture, connectors, and trust posture

Logical build plan for prototype and MVP evolution

Logical architecture


- Experience: Home, Decision Hub, Execution Hub, brief export.
- API and Auth: single gateway, RBAC, audit logs.
- Orchestration: scheduler plus workflow engine for signal refresh and decision updates.
- Connector services: ERP, CRM, HR, email, market, news adapters.
- AI/ML core: vector search, patterning services, model gateway.
- Data layer: signal store, document store, cache, queues.

Connectors and trust posture

- Treat BI outputs as signal sources, not as the product itself.
- Tier 0: CSV and file uploads; Tier 1: standard connectors; Tier 2: connector SDK and certification.
- Security baseline: tenant isolation, encryption, auditability, data minimization.
- No training on customer data by default, explicit opt-in only.
- Support deployment portability and multiple deployment modes as the product matures.
- Keep implementation cloud-portable even when using AWS as the first reference stack.

A7. AWS reference implementation

Use this as the implementation blueprint, not as a product lock-in statement



How to talk about this slide

Edge + Frontend: S3/CloudFront/WAF for delivery and protection.

API + Auth:

API Gateway + Cognito + Secrets Manager.

Orchestration:

Step Functions + EventBridge for refresh and workflow timing.

Agents:

separate ERP, CRM, HR, Email, Market, News agents.

AI/ML core:

Bedrock, OpenSearch, SageMaker as one implementation path.

Data + Messaging:

Aurora, DynamoDB, S3, Redis, SQS/SNS.

Positioning:

reference implementation on AWS, but product remains cloud-portable.

A8. Developer checklist before the first prototype review

Use this slide as the acceptance gate before showing the prototype again

Developer checklist

- Home screen shows Top 3 decisions, What changed, and Value at stake with mock but structured data.
- Decision Hub exposes Why now, Impact range, Confidence, Top 3 evidence, Risks, and Commit/Hold.
- Execution Hub tracks decision, next action, owner, due date, status, and value-at-stake.
- All visible data objects map back to the canonical Signal schema.
- Evidence and confidence are present in the UI, not hidden in notes.
- Actions created from decisions persist across screens and are reflected in Execution Hub.
- No dead-end clicks. Every primary action either drills down, commits, exports, or updates status.

A8. Data set

-

Dummy Company

A) Company identity and operating shape

1. What should we call the company, and do you want a ticker symbol shown (yes/no)?
 - a. **Gaksh Industrial Systems (GIS)**
2. Fiscal year end: Dec or Mar or other?
 - a. **December**
3. Main operating units: do you want to show it as 3 regions (Americas, EMEA, APAC) or by countries (US, UK, India)?:
 - a. **US, UK, India**

B) Footprint details

4. What are the key sites you want visible in the prototype? Example: Plano plant (US), Pune plant (India), Birmingham (UK). Any more?
 - a. **Detroit plant (US), Detroit plant (US) ,Pune plant (India), Michigan (UK)**
5. For each site, what is its role: manufacturing, R&D, sales, distribution, shared services?
 - a. **All: R&D, sales, distribution, shared services, Pune also has shared services center supporting global IT and customer services**

C) Products and customer segmentation

6. Top 2–3 product lines you want in filters: Car Axle, braking systems, steering, electronics, powertrain, etc?
 - a. **Car Axle, braking systems, steering**
7. Customer segments: only OEM, or OEM + Tier-1 + aftermarket?
 - a. **OEM + Tier-1**
8. Name 3–5 “Customer A/B/C” OEMs you want to appear (real names ok, or anonymized)?
 - a. **Ford Motors (US), Bajaj Auto (India) , Jaguar (UK)**

D) Commercial model (for pipeline and decisions)

9. Typical deal types: multi-year supply contracts, annual renegotiations, spot buys, services, retrofits?
 - a. **multi-year supply contracts, annual renegotiations**
10. Average deal size bands you want to show: for example \$3–8M per deal, or larger?
 - a. **1 M USD smallest deal, 3 M mid range deals , 8 M USD large deals**

E) Vendors and supply chain reality

11. Top 3 supplier regions that matter (China, Mexico, Eastern Europe, etc)?
 - a. **India, China, US**
12. Any “critical component” examples to use for D1: bearings, forged steel, sensors, castings?
 - a. **Sensors and castings**

F) KPIs that appear on the Home page are below

- Order intake / bookings (30D and QTD)
- Backlog and book-to-bill
- Operating margin
- Free cash flow
- OTIF (on-time in-full)
- Safety LTIFR

G) Decision set for the demo

- **D1: Dual-source critical component**
- **D2: Quality containment for Product Line B**
- **D3: Services + retrofit growth play**
- **D4: plant rebalancing)**

H) Execution realism

13. Who owns execution in your org for the demo: COO, CRO, Head of Procurement, Head of Quality, Plant Manager?
 - a. **Execution is owned by COO, COO reports to white CEO who in turn reports to the board**
14. What statuses do you want: Not started, In progress, Blocked, Done, Hold?
 - a. **Not started, In progress, Blocked, Done, Hold?**

SiBoNi GIS v3 Dataset

Developer Guide (Prototype + MVP Scaffolding)

Core + Optional tables, screen-to-table map, API examples

Simple rule to prevent over complication while making Prototype

- Dataset can have 15–20 tables.
- Prototype UI should only “touch” 6–8 tables.

What this dataset is

How to use the v3 synthetic dataset to build the SiBoNi prototype now and keep it MVP-ready.

CORE

OPTIONAL

- Purpose: seed data for end-to-end flows (Home → Intelligence → Decision → Execution).
- Contents: CSV tables + API example payloads (JSON).
 - Use `/api_examples` for immediate UI wiring.
 - Swap JSON reads with DB/API calls later without changing screen logic.
- Two tiers:
 - `CORE_*.csv` = build against these now (prototype + MVP scaffolding).
 - `OPTIONAL_*.csv` = included for later (do not wire into UI yet unless needed).
- Reference maps:
 - `TABLES_MANIFEST.csv` (table list + tier).
 - `SCREEN_TO_TABLE_MAP.csv` (which screens use which tables).
- Every row includes `tenant_id` (TENANT-GIS) for multi-tenant readiness.

Quick start (fastest path)

Recommended approach to get a working prototype in the shortest time.

- Unzip dataset to a local folder (e.g., /data/gis_v3).
- Wire the UI to JSON payloads first (no backend required):
 - `api_examples/home_payload.json` → Home screen
 - `api_examples/decision_pack_DR-001..DR-004.json` → Decision Hub drill-down
- Build list/table screens against the CSVs in parallel:
 - Signals table → `CORE_13_signals_last90d.csv`
 - Execution table → `CORE_16_actions_execution_D1_D4.csv`
- When UI stabilizes, replace JSON reads with DB/API queries using the same shapes.
- Keep OPTIONAL tables stored but not surfaced in the prototype UX.

Data conventions (how to join)

Key IDs and relationship tables you will use throughout the app.

- Tenant model:
 - `tenant_id` is present on every table (TENANT-GIS in this dataset).
- Primary join keys:
 - Operational slice: `site_id`, `product_line_id`, `segment_id`, `country`, `date`
 - Commercial slice: `customer_id`, `deal_id`
 - Decision graph: `decision_id`, `signal_id`, `action_id`, `initiative_id`
 - Supply chain: `supplier_id`, `component_id`
- Relationship tables:
 - Decision ↔ Signals: `CORE_15_decision_evidence.csv`
 - Decision ↔ Initiatives: `CORE_17b_initiative_alignment.csv`
 - Actions ↔ Status history: `CORE_16b_action_updates.csv`
- All timestamps are ISO 8601; dataset anchor date is 2026-03-14.

Home screen: Decision Agenda

How to populate “What changed”, “Top decision this week”, and KPIs.

- What changed (3 bullets):
 - Query CORE_13_signals_last90d by filter slice (country/site/product/segment).
 - Sort by date desc + confidence desc; take top 3 summaries.
- Top decision this week:
 - From CORE_14_decisions_D1_D4; optionally re-rank using your scoring.
 - Pull: title, why_now, impact_range, confidence, time_window_days.
- KPI rollups (last 30 days):
 - Aggregate CORE_10_kpi_daily_last90d_by_segment for selected slice.
 - Show: Order intake, OTIF, Operating margin, LTIFR (+ backlog/b2b from weekly table if desired).
- Push strip / nudges (optional on Home):
 - CORE_19_alerts filtered by user_id; show New/Seen.

Intelligence Hub (Signals table)

Primary data sources and recommended filters.

- Primary table:
 - CORE_13_signals_last90d (90 days, all categories).
- Default filters:
 - date range (last 30/60/90), category, country/site, product line, segment
 - confidence threshold, needs_verification flag
- Context strip (“live” indicators):
 - CORE_22_external_indicators + CORE_23_external_indicator_values_last90d
- Commercial context (optional):
 - CORE_12_deals_pipeline for late-stage and high-value activity.

Decision Hub (Decision brief)

Tables needed to show options, evidence, assumptions, and alignment.

- Decision record:
 - CORE_14_decisions_D1_D4 (why_now, impact_range, confidence, time_window).
- Options:
 - CORE_14b_decision_options (recommended vs conservative).
- Evidence pack:
 - Join CORE_15_decision_evidence → CORE_13_signals_last90d.
 - Display Top 3 evidence bullets with provenance + confidence.
- Decision rigor blocks:
 - Assumptions: CORE_14c_decision_assumptions
 - Leading indicators: CORE_14d_decision_leading_indicators
 - Pre-mortem: CORE_14e_decision_premortem
- Alignment and execution:

Execution Hub (Commitments)

Single view: decision → actions → owner → status → due date (+ blockers).

- Primary table:
 - CORE_16_actions_execution_D1_D4 (title, owner, status, due_date, blocker, value_at_stake).
- Status history / timeline:
 - CORE_16b_action_updates (updates by date).
- Initiative overlay (optional, but recommended):
 - CORE_17_initiatives to show broader program context.
- Decision lifecycle (optional):
 - CORE_18_decision_events for “committed/held/closed” trail.
- Use alerts (CORE_19_alerts) to trigger reminders for blocked / due-soon actions.

CORE vs OPTIONAL tables

How to avoid prototype bloat while keeping MVP-ready scaffolding.

- Build prototype UX against CORE tables only.
- CORE covers: signals → decisions → commit/hold → actions → progress.
- OPTIONAL tables are included for later readiness:
 - OPTIONAL_01_audit_log (compliance + analytics)
 - OPTIONAL_02_evidence_objects (attachments / URIs / hashes)
 - OPTIONAL_03_supplier_risk_events (event stream)
- Rule of thumb:
 - Dataset can have many tables; prototype UI should “touch” only 6–8 core domains.

Demo realism tips

How to keep “last 30 days” and “last updated” consistent in the demo.

- Last updated:
 - Use $\max(\text{date})$ from signals/KPIs as the “Last updated” timestamp in the UI.
- Last 30 days filter:
 - Compute $\text{start} = \max(\text{date}) - 29 \text{ days}$; filter by $\text{date} \geq \text{start}$.
- Time shifting for repeat demos:
 - If you want a “moving story”, shift all date fields by a fixed offset in a preprocessing step.
- Keep the narrative tight: What changed → Top decision → Commit → Execution.

Recommended minimal API endpoints

Thin backend contract that matches the wireframe screens.

- GET /home?country=&site_id=&product_line_id=&segment_id=&window=30d
 - Returns: what_changed[], top_decision, value_at_stake, kpi_rollups, alerts[]
- GET /signals?filters...
 - Returns paginated signals with provenance, confidence, needs_verification.
- GET /decisions?filters...
 - Returns ranked decisions list; includes confidence, impact_range, time_window.
- GET /decisions/{id}
 - Returns decision pack: options, evidence, assumptions, indicators, premortem, alignment, actions.
- GET /execution?filters...
 - Returns actions table + updates; supports status and due-date filters.
- POST /decisions/{id}/commit or /hold
 - Creates decision_event + actions; triggers alerts.

Key files to open first

When you unzip the dataset, start here.

- TABLES_MANIFEST.csv
 - Authoritative list of tables and which tier they belong to.
- SCREEN_TO_TABLE_MAP.csv
 - What each screen should query.
- api_examples/home_payload.json
 - Use this to wire Home immediately.
- api_examples/decision_pack_DR-001.json
 - Use this to wire Decision Hub drill-down immediately.
- CORE_13_signals_last90d.csv (Signals table) and CORE_16_actions_execution_D1_D4.csv (Execution table).